

Module 6

Statistics and indexes

Copyright Tibor Karaszi Konsulting and Cornerstone Group AB

Module Overview

- Statistics and cardinality estimation
- Indexes
- Columnstore indexes

Lesson: Statistics and cardinality estimation

- Cost-based optimization
- Predicate selectivity
- Inspecting statistics
- Cardinality estimation
- Optimizer fixes
- Optimizer enhancements
- Controlling CE, OF and OE
- Creating statistics
- Updating statistics
- Filtered statistics

Cost-based optimization

- Cost-based optimization selects the query execution plan with the lowest estimated cost for execution
- Statistics are a critical part of cost-based optimization
 - Statistics provide information about data distribution in a column or group of columns
 - Columns in tables and indexes might have statistics

What Is a Cost-Based Optimizer?

<https://www.brentozar.com/archive/2021/09/what-is-a-cost-based-optimizer/>

Predicate selectivity

- Predicates
 - Expressions which evaluate to true or false
 - Found in joins, WHERE and HAVING clauses
- Predicate Selectivity
 - The number of rows from a table that meet a predicate
 - High selectivity = small percentage of rows returned
 - Low selectivity = large percentage of rows returned
- Predicate Selectivity in Query Optimization
 - Selectivity used to sequence join operations and choose between scan/seek

Query Tuning Fundamentals: Density, Predicates, Selectivity, and Cardinality

<https://sqlserverdbknowledge.wordpress.com/2014/02/24/query-tuning-fundamentals-density-predicates-selectivity-and-cardinality/>

Inspecting statistics

- Statistics header
 - Statistics metadata
- Density vector
 - Measure of uniqueness
- Histogram
 - Data distribution, up to 200 steps
- You can investigate statistics using DBCC SHOW_STATISTICS

Demo Show statistics

Statistics

<https://docs.microsoft.com/en-us/sql/relational-databases/statistics/statistics>

What are SQL Server Statistics and Where are they Stored?

<https://www.sqlnethub.com/blog/what-are-sql-server-statistics-and-where-are-they-stored/>

DBCC SHOW_STATISTICS (Transact-SQL)

<https://docs.microsoft.com/en-us/sql/t-sql/database-console-commands/dbcc-show-statistics-transact-sql>

Cardinality estimation

- In SQL Server, cardinality estimation attempts to predict the number of rows returned by a query or query operator
- SQL Server 2014 and SQL Server 2016 include rewritten cardinality estimation logic
 - Controlled by database compatibility level
 - Minor enhancements in subsequent releases
- Several factors can cause poor cardinality estimation:
 - Out-of-date or missing statistics
 - Functions in predicates
 - Table variables

Cardinality Estimation (SQL Server)

<https://docs.microsoft.com/en-us/sql/relational-databases/performance/cardinality-estimation-sql-server>

Compatibility Levels and Cardinality Estimation Primer

<https://sqlperformance.com/2019/01/sql-performance/compatibility-levels-and-cardinality-estimation-primer>

Optimizer fixes

- A weakness in the optimizer can be fixed
- Prior to 2016 we had to opt in:
 - Using the trace flag for that particular fix
 - Or 4199 which enables all fixes
- As of 2016 the fixes of RTM are enabled
 - Assuming 2016 compatibility level
 - Patching SQL server will not add fixed
 - Unless you use trace flag 4199

Enabling SQL Server Optimizer Hotfixes

<https://sqlrus.com/2020/07/understanding-how-to-enable-sql-server-optimizer-hotfixes/>

Optimizer enhancements

- New functionality is added to the optimizer
- One example is "Batchmode over rowstore"
 - Added in 2019
- Enabled by the database compatibility level
 - Same as when enhancement was released
 - Or higher

Intelligent query processing in SQL databases

<https://docs.microsoft.com/en-us/sql/relational-databases/performance/intelligent-query-processing>

SQL SERVER 2022 Features for Performance Optimization

<https://blog.sqlauthority.com/2021/11/08/sql-server-2022-features-for-performance-optimization/>

Controlling CE, OF and OE

Setting	Type	CE	OF	OE
Db compat level	DB	Y	Y	Y
QUERY_OPTIMIZER_COMPATIBILITY_LEVEL_xxx	H	N	Y	Y
QUERY_OPTIMIZER_HOTFIXES	DB	N	Y	N
ENABLE_QUERY_OPTIMIZER_HOTFIXES	H	N	Y	N
FORCE_LEGACY_CARDINALITY_ESTIMATION	H	Y	N	N
LEGACY_CARDINALITY_ESTIMATION	DB	Y	N	N
9481 (force legacy CE)	TF	Y	N	N
2312 (force current CE with db compat 2012)	TF	Y	N	N

DB: Database setting

H: Hint

TF: Trace flag, can also be turned on at query level with
QUERYTRACEFLAG hint

Demo CE and optimizer fixes

Creating statistics

- Automatic creation:
 - When AUTO_CREATE_STATS = ON
 - Single column statistics only
 - Names start _WA...
- Manual creation:
 - CREATE STATISTICS

Create Statistics

<https://docs.microsoft.com/en-us/sql/relational-databases/statistics/create-statistics>

Updating statistics

- Automatic update:
 - AUTO_UPDATE_STATISTICS = ON
 - Automatic update threshold
 - AUTO_UPDATE_STATS_ASYNC option
- Manual update:
 - UPDATE STATISTICS command
 - sp_updatestats
- Updating statistics causes query plans to be recompiled
- Database maintenance:
 - Maintenance plans will update statistics all statistics
 - Even if we have modified 0 rows since last time – waste of resources
 - Ola Hallengren's IndexOptimize focuses on defragmentation
 - Will only update statistics if it happens to do a rebuild on the index
 - Create your own job that calls IndexOptimize with parameters to update statistics only

Demo Update statistics

SQL Server Statistics and how to perform Update Statistics in SQL

<https://www.sqlshack.com/sql-server-statistics-and-how-to-perform-update-statistics-in-sql/>

UPDATE STATISTICS (Transact-SQL)

<https://docs.microsoft.com/en-us/sql/t-sql/statements/update-statistics-transact-sql>

Update Statistics

<https://docs.microsoft.com/en-us/sql/relational-databases/statistics/update-statistics>

Filtered statistics

- CREATE STATISTICS using a WHERE clause
- Some limitations:
 - AUTO_UPDATE_STATISTICS threshold is based on all rows in table
 - Not the rows that the filter specifies
 - Limited to simple comparison logic
 - Cannot be created on all data/column types
 - Cannot be created on indexed views
- Histogram has finer resolution
 - When you combine several statistics

Filtered Statistics Follow-up

<https://www.brentozar.com/archive/2016/12/filtered-statistics-follow/>

CREATE STATISTICS (Transact-SQL)


<https://docs.microsoft.com/en-us/sql/t-sql/statements/create-statistics-transact-sql>

Lesson: Indexes

- Heap internals
- Index structure
- Picking the right index key
- Single column and multicolumn indexes
- Filtered indexes
- The query optimizer's choice of indexes
- Data modification internals
- Index fragmentation
- Index column order
- Identify and create missing indexes

Heap internals

- A table without a clustered index is a heap
- IAM pages contain pointers to extents containing heap data
 - The IAM is the only link between pages in a heap; row data is unordered
- Heap rows are identified by a row identifier (RID)
 - Nonclustered indexes on a heap contain RID pointers



Heaps (Tables without Clustered Indexes)

<https://docs.microsoft.com/en-us/sql/relational-databases/indexes/heaps-tables-without-clustered-indexes>

Tables Without Clustered Indexes

<https://www.brentozar.com/blitz/heaps-tables-without-primary-key-clustered-index/>

Heaps in SQL Server: Part 1 The Basics


<https://www.red-gate.com/simple-talk/databases/sql-server/t-sql-programming-sql-server/heaps-in-sql-server-part-1-the-basics/>

Clustered Index vs. Heap in SQL Server

<https://www.sqlshack.com/clustered-index-vs-heap/>

Index structure

- B-trees:
 - Self-balancing tree data structure
 - One root node; many non-leaf nodes; many leaf nodes
 - Clustered and nonclustered indexes are b-trees
- Index key:
 - Defines the order of data in an index
- Clustered index:
 - Leaf node *is* the data, contains all columns
 - Higher levels has index key
 - Index order = table data order
- Nonclustered index:
 - All nodes contain index pages
- There are other index types than B-Trees:
 - XML, full-text, hash, columnstore and spatial indexes



Indexes

<https://docs.microsoft.com/en-us/sql/relational-databases/indexes/indexes>

Clustered and nonclustered indexes described

<https://docs.microsoft.com/en-us/sql/relational-databases/indexes/clustered-and-nonclustered-indexes-described>

Picking the right index key

- Clustered index key criteria:
 - Unique
 - Non-nullable
 - Narrow
 - Static
 - Ever-increasing
- Nonclustered index criteria:
 - Frequently used predicates
 - Join columns
 - Aggregate queries
 - Avoid:
 - Redundant/duplicate indexes
 - Wide keys
 - Indexes serving only one query for systems with a lot of data change

Designing effective SQL Server clustered indexes

<https://www.sqlshack.com/designing-effective-sql-server-clustered-indexes/>

Clustered and Non Clustered Index: 7 Top Points Explained

<https://codingsight.com/clustered-and-non-clustered-index-7-top-points-explained/>

Single column and multicolumn indexes

- Single column index:
 - Each predicate column has its own index
 - Less performance gain, but more reusable
- Multicolumn index:
 - All predicate columns are included in the key of the index
 - Greatest performance gain, but limited reusability

Demo Index several columns

Filtered indexes

- Nonclustered index with a filter predicate:
 - Filter predicate defined with a WHERE clause in the index definition
- Better performance than a whole-table index:
 - Finer-grained statistics
 - Smaller size
- Suitable when table data is queried in clearly identifiable subsets
- Some restrictions apply

Demo Filtered index

Create filtered indexes

<https://docs.microsoft.com/en-us/sql/relational-databases/indexes/create-filtered-indexes>

What You Can (and Can't) Do With Filtered Indexes

<https://www.brentozar.com/archive/2013/11/what-you-can-and-cant-do-with-filtered-indexes/>

The query optimizer's choice of indexes

- Predicate SARGability:
 - WHERE <column> <operator> <value>
 - <column> must appear alone
 - Leading string wildcards prevent index seek

Demo SARG

Sargable

<https://en.wikipedia.org/wiki/Sargable>

The Two Ways to Fix Non-Sargable Queries

<https://www.brentozar.com/archive/2019/12/the-two-ways-to-fix-non-sargable-queries/>

How to use sargable expressions in T-SQL queries; performance advantages and examples

<https://www.sqlshack.com/how-to-use-sargable-expressions-in-t-sql-queries-performance-advantages-and-examples/>

Data modification internals

- Delete:
 - Key reference removed from leaf-level page
 - A page with no references is removed from the b-tree
- Insert:
 - New key reference added to either:
 - Existing free space on a page
 - New page
- Update:
 - A delete followed by an insert

Index fragmentation

- External fragmentation
 - Linked list jumps back and forth in the database file
 - Less relevant unless you have single magnetic disks
- Internal fragmentation
 - Pages are less than 100% full
 - Watch out for
 - Rebuilding and re-applying a lower fill-factor
 - Rebuilding and **not** re-applying a lower fill-factor



Demo Fragmentation

Optimize index maintenance to improve query performance and reduce resource consumption
<https://docs.microsoft.com/en-us/sql/relational-databases/indexes/reorganize-and-rebuild-indexes>

Index fragmentation revisited

<https://sqlblog.karaszi.com/index-fragmentation-revisited/>

Stop Worrying About SQL Server Fragmentation

<https://www.brentozar.com/archive/2012/08/sql-server-index-fragmentation/>

Rebuild all fragmented heaps

<https://karaszi.com/rebuild-all-fragmented-heaps>

Index column order

- SELECT performance:
 - Only the first column has a histogram
 - Use the most selective column as the first column, but consider how the index will be used
- INSERT performance:
 - Consider the effect column order will have on INSERT performance for a clustered multicolumn index

How to Think Like the Engine: Index Column Order Matters a LOT.

<https://www.brentozar.com/archive/2019/11/how-to-think-like-the-engine-index-column-order-matters-a-lot/>

Identify and create missing indexes

- Query plans
- Query store
- Database Engine Tuning Advisor
- Missing index DMVs

- Tools will not:
 - Suggest clustered indexes
 - Suggest modifications to existing indexes
 - Analyze column order in multicolumn indexes
 - Suggest filtered indexes

Demo Missing index views

Index Related Dynamic Management Views and Functions (Transact-SQL)

<https://docs.microsoft.com/en-us/sql/relational-databases/system-dynamic-management-views/index-related-dynamic-management-views-and-functions-transact-sql>

Diagnosis: Index-a-phobia

https://www.brentozar.com/blitzindex/sp_blitzindex-missing-indexes/

Glenn Berry's SQL Server Diagnostic Queries

<https://glennsqlperformance.com/resources/>

Lesson: Columnstore indexes

- Columnstore indexes
- Columnstore index features
- Columnstore index recommendations
- Maintaining Columnstore indexes

Columnstore indexes

Rowgroups (1 million rows per group)

Segments (one per column)

A	B	C
14	88	57
65	3	78
12	76	2
31	12	87
45	89	46
19	22	45
8	76	89

Segments (one per column)

A	B	C
11	18	5
65	99	44
52	2	65
71	98	8
5	8	32
43	12	66
4	8	45

Segments (one per column)

A	B	C
33	36	44
87	54	6
4	82	77
12	13	36
65	9	94
28	8	24
1	65	3

...WHERE B = 91

Deltastore (open)

Delete bitmap

Demo Columnstore

Columnstore indexes: Overview

<https://learn.microsoft.com/en-us/sql/relational-databases/indexes/columnstore-indexes-overview>

What are Columnstore Indexes?

<https://www.red-gate.com/simple-talk/databases/sql-server/t-sql-programming-sql-server/what-are-columnstore-indexes/>

Columnstore Index Scan

<https://sqlserverfast.com/epr/columnstore-index-scan/>

Niko Neugebauer, Columnstore

<https://www.nikoport.com/columnstore/>

Columnstore index features

- Clustered or nonclustered
- We can have both row and column indexes on the same table
 - Only one columnstore index
- Columnstore index options
 - Filter
 - Compression delay
- Data is highly compressed
 - Inserts and after-image for update are initially reflected in row-based delta-store
 - After 1 million rows, such open group is compressed
 - Delete and before-image of update are reflected in delete bitmap
- ORDER clause for clustered Columnstore index (2022)
 - Data is potentially sorted when the index is built
 - Data is sorted (per batch) when data is added

CREATE COLUMNSTORE INDEX (Transact-SQL)

<https://docs.microsoft.com/en-us/sql/t-sql/statements/create-columnstore-index-transact-sql>

Columnstore indexes - what's new

<https://docs.microsoft.com/en-us/sql/relational-databases/indexes/columnstore-indexes-what-s-new>

Stairway to Columnstore Indexes

<https://www.sqlservercentral.com/stairways/stairway-to-columnstore-indexes>

Columnstore index recommendations

- Clustered Columnstore indexes
 - Fact tables
 - Large dimensional tables
 - Large log tables (typically append-only, potentially huge)
- Non-clustered Columnstore indexes
 - OLTP system where you also perform analysis
 - Not uncommon that many indexes has been added over time
 - Indexes added for analysis queries can potentially be replaced with:
 - One (1) Columnstore index, vastly smaller in size than the b-tree indexes combined
- Do not remove b-tree indexes created for high selectivity queries
 - No SEEKs in a Columnstore index

Columnstore indexes - what's new

<https://docs.microsoft.com/en-us/sql/relational-databases/indexes/columnstore-indexes-what-s-new>

Niko Neugebauer, Columnstore

<https://www.nikoport.com/columnstore/>

Maintaining Columnstore indexes

- What is the purpose of the index?
 - Compression only
 - Used for queries, where you also expect potential segment elimination
- What modifications are performed?
 - If data was removed (DELETE or UPDATE)
 - Consider getting rid of the delete bitmap
 - If data was added (INSERT or UPDATE)
 - And you expect segment alignment
 - ...then you want to re-process the index
 - ALTER INDEX REBUILD of an ORDERed clustered Columnstore index
 - Soft sort, which doesn't re-align all data
 - Consider converting to a row-store and then back again to Columnstore
 - Using CREATE INDEX ... WITH DROP_EXISTING

Demo Columnstore ordered

ORDER for columnstore index

<https://sqlblog.karaszi.com/order-for-columnstore-index/>

Performance tuning with ordered clustered columnstore index

<https://learn.microsoft.com/en-us/azure/synapse-analytics/sql-data-warehouse/performance-tuning-ordered-cci>

Columnstore indexes - Data loading guidance

<https://learn.microsoft.com/en-us/sql/relational-databases/indexes/columnstore-indexes-data-loading-guidance>

Lab 6: Statistics and indexes

- Ex 1: Improve a query, the easy way
- Ex 2: Improve a query, the traditional way
- Ex 3: Improve a query, by query re-write

Estimated Time: 30 minutes